



FULLY-DYNAMIC ALL-PAIRS OF SHORTEST PATHS

Fully Dynamic(=Online) APSP

- N vertices
- Directed
- weighted(no negative cycles) or unweighted
- Begins with empty graph
- Inserts or delete vertices in each update
- Answers queries of the form “distance between shortest path from s to t ” in $O(1)$
- (optional but usually possible) Find first k edges of a shortest path in $O(k)$

Randomization

- I will skip all parts about randomization because too much...

- The use of any kind of **randomness** is prohibited. Any solution seen using randomness (whether it provably works or not) will be disqualified. This includes writing your own pseudorandom generators instead of using the `rand()` function.

Previous Results

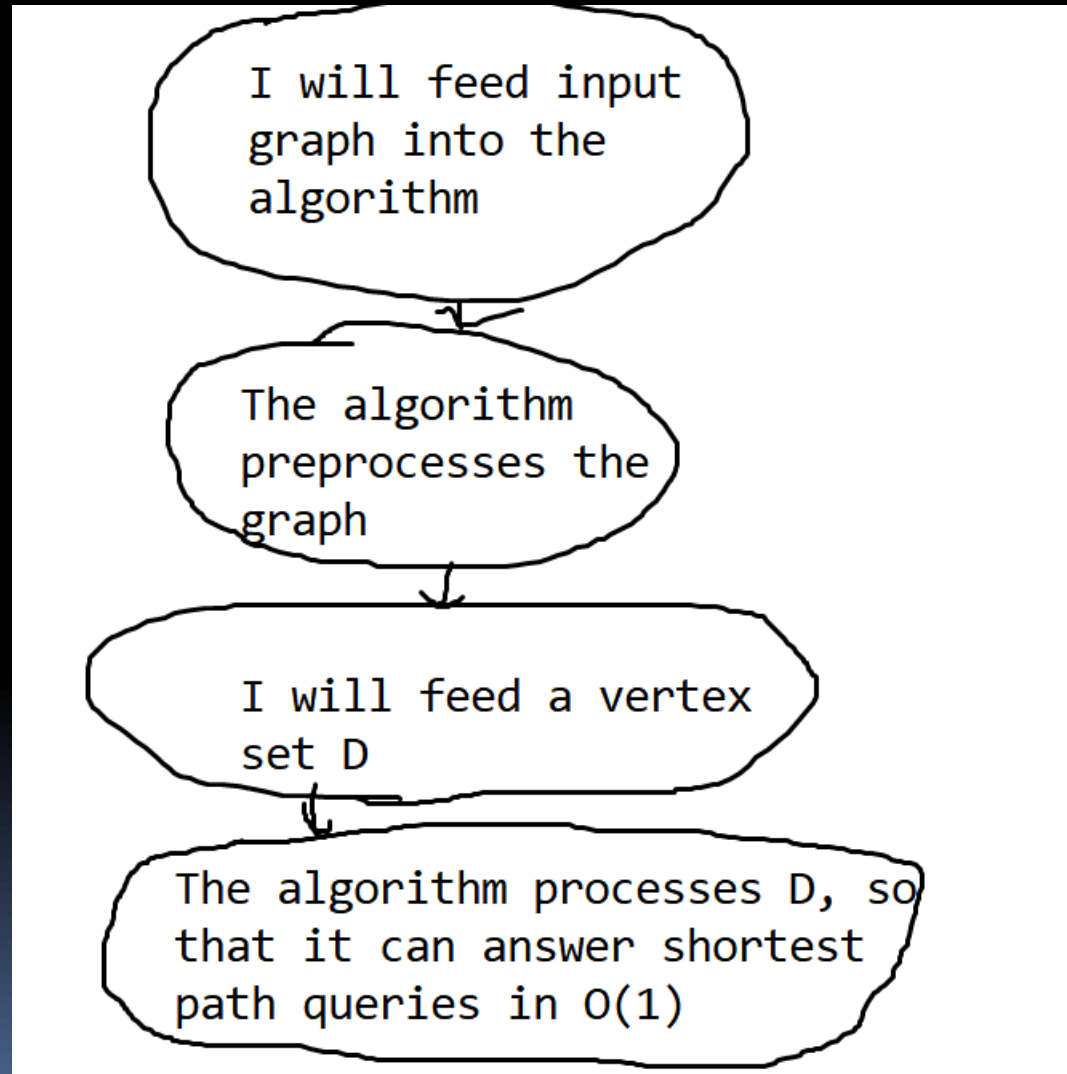
	Amortized update time	Worst case update time	Weighted	Space
Brute	$O(n^3)$	$O(n^3)$	Yes	$O(n^2)$
Dlo ₄ ,THOo ₄	$O(n^2 \log^2 n)$	$O(n^3)$	Yes	$O(n^3)$
THOo ₅	$\tilde{O}(n^{2+3/4})$	$\tilde{O}(n^{2+3/4})$	w>=0	Super-cubic
Gutenberg, Nilsen 2020	$O(n^{\frac{19}{7}} \log^{8/7} n)$	$O(n^{\frac{19}{7}} \log^{8/7} n)$	Yes	Sub-cubic
Gutenberg, Nilsen 2020	$O(n^{2.6} \log n)$	$O(n^{2.6} \log n)$	No	Sub-cubic
Gutenberg, Nilsen 2020	$O(n^{\frac{11}{4}} \log^{2/3} n)$	$O(n^{\frac{11}{4}} \log^{2/3} n)$	Yes	$O(n^2)$
Gutenberg, Nilsen 2020	$O(n^{\frac{8}{3}} \log^{2/3} n)$	$O(n^{\frac{8}{3}} \log^{2/3} n)$	No	$O(n^2)$

Definitions

- h -hop path: Path with at most h edges (not same as distance)
- $hop(p)$: number of edges on path p
- $dist_H(s, t)$: shortest path from s to t in the induced subgraph regarding H
- Improving path from s to t : path from s to t with distance at most $dist(s, t)$
- Improving path from s to t with regards to H : path from s to t (the path doesn't have to be in H) with distance at most $dist_H(s, t)$

Reduction to batch deletion

- “Batch Deletion Problem”



Reduction to batch deletion

- Imagine we have a data structure (call it P) that:
 - Inputs a graph and preprocesses in time $O(t_{pre})$
 - Can handle one single batch deletion of a vertex set D with size $\leq 2B$ in $O(t_{del})$
 - Returns shortest path between all pairs of nodes in new graph (after deletion)
- We can use it to solve the fully dynamic APSP in worst update time $O\left(\frac{t_{pre}}{B} + t_{del} + Bn^2\right)$
- Standard deamortization techniques

Reduction to batch deletion

- Insertion = Easy
- Modified Floyd-Warshall
 - Compute distance from each original node to new node
 - Compute distance from new node to each original node
 - Update all $dist(s, t)$ with $dist(s, new) + dist(new, t)$
- $O(kn^2)$ for inserting k vertices

Reduction to batch deletion

Assume $B=5$

Queries in between updates

Updates	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
Instance 1																				
Instance 2																				

→ Build P from graph
after updates 1~5

Split the
preprocessing work
throughout updates 6
~10 evenly in
 $O(t_{\text{pre}}/B)$ each

→ Use P, process
updates 6~14

Reduction to batch deletion

- Note that in the 11th update, you answer queries using P and process update 6~11
- Then for the 12th update, you discard the processing and start over using P and process updates 6~12
- A bit confusing, took me a while to understand

Reduction to batch deletion

- The preprocessing time is spread across B updates and therefore worst update time for this step is $O\left(\frac{t_{pre}}{B}\right)$
- When you use P and process updates
 - Process deletions first in $O(t_{del})$
 - Use the results and apply modified Floyd-Warshall to process insertion updates in $O(Bn^2)$
- $O\left(\frac{t_{pre}}{B} + t_{del} + Bn^2\right)$ worst update time total

Powerful Blackbox


- Given graph G
- Given a parameter h
- Given n^2 paths $p_{s,t}$ such that
 - ▣ If shortest path (if many, the one with least edges) from s to t in G consists of at most h edges (h -hop), then $p_{s,t}$ is this path.
- Returns all pairs of shortest path in G
- Time Complexity: $O\left(\frac{n^3 \log n}{h} + n^2 \log^2 n + n^2 h\right)$

Powerful Blackbox

- In other words
- If we already found all h -hop improving shortest paths with regards to G
 - That is, if there exist path p such that
 - Starts at s and ends at t
 - $hop(p) \leq h$
 - $dist(p) = dist(s, t)$
 - Then $p_{s,t}$ will be one of the valid p
 - Otherwise $p_{s,t}$ can store any valid path between s, t
- We can find all pairs of shortest path



Powerful Blackbox

- Will describe near the end of the presentation
 - Lets find all h -hop improving shortest paths first
- 

Slow Deletion

- Motivation

- If we only maintain shortest paths that consist of little edges
- Handling deletion is easier because we only need to recompute paths that includes at least one of the elements in the deleted set
- As the paths are short, the probability of recomputation is small
- We can then use our powerful blackbox to extend our short paths to all pairs of shortest path

Slow Deletion


- Objective: Find all h -hop improving paths in $G \setminus D$
 - D is the set to be deleted
 - Well, not exactly “all”, we only need to find one path between each pair of vertices
 - Very important to understand what it means
 - If the shortest path from s to t that consist of least edges has $\leq h$ edges, it is an h -hop improving path
 - We don't have to know which are h -hop improving paths and which aren't, we just have to make sure that every such path in our set of paths
- Method: Recursion/Induction

Slow Deletion

- Lets say we are finding all h -hop improving paths that starts from s
- Imagine we have all $\frac{2h}{3}$ -hop improving paths between all pairs of vertices already
- Pick some integer $r \in [\frac{h}{3}, \frac{2h}{3}]$ and let Sep be the set of all vertices x such that the shortest path (which is already found) from s to x consists of exactly r edges

Slow Deletion

Case 1




some h -hop improving path

This path is also a $2h/3$ -hop improving path

Found in previous stage already

Case 2



both r and $\text{length} - r$ are $\leq 2h/3$

$s \rightsquigarrow x$ and $x \rightsquigarrow t$ are $2h/3$ -hop improving paths

Slow Deletion

- So, to find the h -hop improving path from s to t
- We will enumerate x from Sep , and see if the concatenation of $s \leadsto x$ and $x \leadsto t$ is shorter than the current path we have from s to t
- As there is $O(h)$ choices for r , and there are at most $O(n)$ paths we have starting from s
- We can pick r such that $|Sep| = O(\frac{n}{h})$

Slow Deletion

- Some Notations for convenience
 - $i_h = \lceil \log_{1.5} h \rceil$ is the number of phases of “binary lifting” we will do (in fact it is 1.5-lifting)
 - $1 \leq i \leq i_h$ is the phase number of the current phase
 - $h_i = 1.5^i$ is the target hop number that you want to find all shortest path in $V \setminus D$ of hops $\leq h_i$, during phase i

Slow Deletion

```

1  path p[log(N)/log(1.5)][N][N];
2  int adj[N][N]; //adjacency matrix
3  int cnt[N];
4  bool better(path x,path y){
5      if(dist(x)!=dist(y)) return dist(x)<dist(y);
6      return hop(x)<hop(y);
7  }
8  void slowDelete(vector<int>D, int h){
9      for(s in V\D)
10         for(t in V\D)
11             p[0][s][t]=adj[s][t];
12     i_h=ceil(log(h)/log(1.5));
13     for(int i=1; i<=i_h ;i++){
14         int h_i=pow(1.5,i); //in this iteration, we are computing h_i-hop shortest paths
15         for(s in V\D){
16             int r;vector<int>Sep;
17             //finding the best r
18             /*initialize cnt*/
19             for(x in V\D) cnt[hop(p[i-1][s][x])];
20             r=h_i/3;
21             for(int j in [h_i/3,2*h_i/3]) if(cnt[j]<cnt[r]) r=j;
22             for(x in V\D) if(hop(p[i-1][s][x])==r) Sep.push_back(x);
23         }
24         for(t in V\D){
25             p[i][s][t]=p[i-1][s][t];
26             for(x in Sep){
27                 if(better(p[i-1][s][x]+p[i-1][x][t],p[i][s][t])) p[i][s][t]=p[i-1][s][x]+p[i-1][x][t];
28             }
29         }
30     }
31 }
32

```

fast

$O(n/h_i)$

Remember the definition of i_h and h_i , will be used frequently later

Less Slow Deletion

- Why previous algorithm slow?
 - $O(n/h_i)$ part is computed many times
 - When h_i is small, will take $O(n^3)$
 - Doesn't use preprocessing
- How to improve?
 - Preprocess all h-hop improving paths in G
 - We can skip $O(\frac{n}{h_i})$ part if $p_{s,t}$ doesn't not contain elements in D already

Less Slow Deletion

```
8 void badPreprocess(int h){
9     i_h=ceil(log(h)/log(1.5));
10    vector<int>X=V;
11    while(!X.empty()){
12        int s=X.back();X.pop_back();
13        for(int i=1; i<=i_h ;i++){
14            int h_i=pow(1.5,i);//in this iteration, we are computing h_i-hop shortest paths
15            bellmanford(s,V,h_i);//finds all shortest path from s to t with hop at most h_i
16            /*stores result in p[i][s][t]*/
17        }
18    }
19 }
```

Bellmanford is $O(n^3 h_i)$

Preprocessing is $O(n^3 h)$ due to geometric sum

$$\sum h_i = 1 + 1.5 + 2.25 + \dots + h = \frac{1.5h - 1}{1.5 - 1} = O(h)$$

Less Slow Deletion

```

20 void lessSlowDelete(vector<int>D, int h){
21     //of course, we don't actually rewrite the content in p[] because we have to use them multiple times
22     //we will make a copy instead
23     for(s in V\D)
24         for(t in V\D)
25             p[0][s][t]=adj[s][t];
26     i_h=ceil(log(h)/log(1.5));
27     for(int i=1; i<=i_h; i++){
28         int h_i=pow(1.5,i); //in this iteration, we are computing h_i-hop shortest paths
29         for(s in V\D){
30             int r; vector<int>Sep;
31             /*compute r, Sep like in slowDelete*/
32             for(t in V\D){
33                 O(h_i) → if(p[i][s][t] and D does not intersect) continue;
34                 //to check this, we need to maintain the intermediate nodes in p[i][s][t], not just only hop and dist
35                 p[i][s][t]=p[i-1][s][t];
36                 O(n/h_i) {
37                     for(x in Sep){
38                         best=t;
39                         if(better(p[i-1][s][x]+p[i-1][x][t], p[i-1][s][best]+p[i-1][best][t])) best=x;
40                     }
41                 }
42                 O(h_i) {
43                     p[i][s][t]=p[i-1][s][best]+p[i-1][best][t];
44                     if(hop(p[i][s][t])>h_i) p[i][s][t]=NULL; //to keep the hop of paths in O(h_i)
45                 }
46             }
47         }
48     }
49 }

```

The $O(h_i)$ parts contributed to $O(n^2 h)$ due to geometric sum, similar to previous slide

Fast Deletion

- Why is previous algorithm slow?
 - It does not improve asymptotically, the $O\left(\frac{n}{h_i}\right)$ parts still is computed many times
- How to improve?
 - We can make use that insertion is fast
 - We don't want some vertex to appear in our $p[][][]$ too frequently, as it would be costly to delete
 - We maintain a set C that appears in h -hop improving shortest path very frequently, and ignore them during preprocessing, then after we finish doing `delete()`, we add the C vertices back using modified Floyd-Warshall

Fast Deletion

Threshold (represented as τ later) is $\geq 2n^2$

```
8  int congestion[N];
9  void Preprocess(int threshold, int h){
10     i_h=ceil(log(h)/log(1.5));
11     vector<int>X=V;
12     vector<int>C;//initially empty
13     while(!X.empty()){
14         int s=X.back();X.pop_back();
15         for(int i=1; i<=i_h ;i++){
16             int h_i=pow(1.5,i);//in this iteration, we are computing h_i-hop shortest paths
17             bellmanford(s,V\C,h_i);
18             //finds all shortest path from s to t with hop at most h_i in the induced subgraph regarding h_i
19             for(t in V\C){
20                 for(vertex v in p[i][s][t]){
21                     congestion[v]+=ceil(n/h_i);
22                 }
23             }
24             for(vertex v not in C){
25                 if(congestion[v]>=threshold) C.push_back(v);
26             }
27             /*stores result in p[i][s][t]*/
28         }
29     }
30 }
```

Each occurrence of vertex v in $p[i]$ will cause $O(n/h_i)$ time in Delete()

Fast Deletion

- So in our process, when vertices occurred too much in the paths we found, we will stop considering it
- This way, the paths we get are h_i -hop improving paths regarding $G \setminus C$
 - $\text{dist}(p[i][s][t])$ is at most the distance from s to t with h_i hop in $G \setminus C$
 - $p[i][s][t]$ may contain vertices from C , but it is ok

Fast Deletion

- $\text{Congestion}(v) \leq 2\tau$ for all vertices
 - After each iteration of bellman-ford, at most $n(h_i)(\frac{n}{h_i}) = n^2$ is added to all vertices in total, and congestion never increases once it exceeds τ
- Sum of $\text{Congestion}(v)$ is $O(n^3 \log h)$
 - After each iteration of bellman-ford, at most $n(h_i)(\frac{n}{h_i}) = n^2$ is added to all vertices in total, and bellman-ford is ran for $O(n \log h)$ times.
- $|\mathcal{C}|$ is $O(\frac{n^3 \log h}{\tau})$
 - Obvious using above results

Fast Deletion

```
31 void Delete(vector<int>D, int h){//basically just lessSlowDelete but V\ (D Union C) instead of V\D
32 //of course, we don't actually rewrite the content in p[] because we have to use them multiple times
33 //we will make a copy instead
34     vector<int>Ban=D union C;
35     for(s in V\Ban)
36         for(t in V\Ban)
37             p[0][s][t]=adj[s][t];
38     i_h=ceil(log(h)/log(1.5));
39     for(int i=1; i<=i_h ;i++){
40         int h_i=pow(1.5,i);//in this iteration, we are computing h_i-hop shortest paths
41         for(s in V\Ban){
42             int r;vector<int>Sep;
43             /*compute r,Sep like in slowDelete*/
44             for(t in V\Ban){
45                 if(p[i][s][t] and D does not intersect) continue;//still D here
46                 //to check this, we need to maintain the intermediate nodes in p[i][s][t], not just only hop and dist
47                 p[i][s][t]=p[i-1][s][t];
48                 /*update p[i][s][t] using Sep*/
49             }
50         }
51     }
52     //////////
53     for(vertex v in C\D) insert(v);//insert in p[i_h][s][t] like floyd warshall
54 }
```

Units of time taken here is bounded by sum of congestion of vertices in D

Fast Deletion

- Basically it is almost the same as `lessSlowDelete`, except that we might miss some paths that pass through elements in C
- In phase i , the algorithm will find all h_i -improving paths in $G \setminus D$, if the path does not contain elements in C
- After the algorithm, we will use modified floyd-warshall to insert elements in C and find all paths that we missed
- Lastly we will plug into the blackbox and get our desired distance matrix

Fast Deletion

- Everytime we need to recompute a path in stage i , it takes $O(\frac{n}{h_i})$ time and also gives $\frac{n}{h_i}$ contribution to the congestion of the vertex that is in D and in the path
- Time complexity of the path recomputation path is $O(|D|\tau) = O(B\tau)$
- Time complexity of the insertion of vertices in C is $O(|C|n^2) = O(\frac{n^5 \log h}{\tau})$

Complexity Analysis

- $O\left(\frac{t_{pre}}{B} + t_{del} + Bn^2\right)$
- $t_{pre} = O(n^3 h)$
- $t_{del} = O\left(B\tau + \frac{n^5 \log h}{\tau}\right) + O\left(\frac{n^3 \log n}{h} + n^2 \log^2 n + n^2 h\right)$
- Choose $\tau = n^{2.25} \log^{0.5} n, h = n^{0.25} \log^{0.5} n, B = n^{0.5}$
- Complexity $O(n^{2.75} \log^{0.5} n)$

Complexity Analysis

- When unweighted, the bellman-ford during preprocessing can be replaced by BFS
- Complexity can be reduced to Complexity $O(n^{\frac{8}{3}} \log^{\frac{2}{3}} n)$
- Choice of parameters is left as exercise to readers (wasn't given in the paper)

Powerful Blackbox

- Given graph G
- Given a parameter h
- Given n^2 paths $p_{s,t}$ such that
 - ▣ If shortest path (if many, the one with least edges) from s to t in G consists of at most h edges (h -hop), then $p_{s,t}$ is this path.
- Returns all pairs of shortest path in G
- Time Complexity: $O(\frac{n^3 \log n}{h} + n^2 \log^2 n + n^2 h)$

Powerful Blackbox

- Intuition




- Previously we had short paths
- This time we have long paths, what now..?
- If we have some set of vertices Sep , such that many paths of a certain length will include at least one vertex from Sep
- As paths are long, maybe we can bound $|Sep|$?

Powerful Blackbox

Lemma A.1 (see [TZ05, RTZ05]). Let $N_1, N_2, \dots, N_n \subseteq U$ be a collection of subsets of U , with $u = |U|$ and $|N_i| \geq s$ for all $i \in [1, n]$. Then, we can implement a procedure $\text{SEPARATOR}(\{N_i\}_{i \in [1, n]})$ that returns a set A of size at most $O(\frac{u \log n}{s})$ with $N_i \cap A \neq \emptyset$ for all i , deterministically in $O(u + \sum_i |N_i|)$ time.

- Example: $u = 5, s = 2, n = 3$

$$U = \{ \text{1}, \text{2}, \text{3}, \text{4}, \text{5} \}$$






Powerful Blackbox

- Proof of lemma: Just repeatedly pick the element that appears in most sets, then delete all sets that have
- In each step n is multiplied by a factor of at most $(1 - s/u)$ (by pigeon-hole principle)
- We know that $\left(1 - \frac{1}{x}\right)^x < 1/e$ for $x \geq 1$
- So after $\frac{u}{s} \lceil \ln n \rceil$ moves, $n < 1$ and thus we obtained Sep of size $O\left(\frac{u \log n}{s}\right)$
- It is very cool because the number of sets only contributes to the size by a log factor


Powerful Blackbox

$$\mathcal{U} = \{ \boxed{1}, 2, 3, 4, 5 \}$$



Powerful Blackbox

$$\mathcal{U} = \{ \boxed{1} \quad \boxed{2} \quad 3 \quad 4 \quad 5 \}$$



Powerful Blackbox

- Back to our problem
- Idea
 - Again, we approach the problem in inductive manner
 - In every step, we extend h by a factor of 1.5, so the idea is that for each path that is than current h but shorter than new h , we want to have some vertices in *Sep* that somewhat is close to the middle of the path
 - We can then run modified floyd-warshall on only vertices in *Sep*
 - We know that for every path that is a shortest path, any subsegment of it is also a shortest path.

Powerful Blackbox

Algorithm 5: DETERMINISTICEXTENDDISTANCES($\Pi = \{\pi_{i_h}(s, t)\}_{s, t}, h$)

Input: A collection of paths Π , that contains a path for each tuple $(s, t) \in V \times V$.

Output: Returns the set of distances $\{(\text{DIST}_{i_{\max}}(s, t))\}_{s, t \in V \times V}$.

```
1 foreach  $(s, t) \in V \times V$  do
2    $\text{DIST}_{i_h}(s, t) \leftarrow w(\pi_{i_h}(s, t))$ 
3 for  $i \leftarrow i_h + 1$  to  $i_{\max}$  do
4   Compute a set SEPARATOR of size  $O(n \log n / h_i)$  that contains a vertex from each
   path in  $\Pi$  of hop at least  $\lfloor \frac{1}{4} h_i \rfloor$ ;
5   foreach  $(s, t) \in V \times V$  do
6      $\text{DIST}_i(s, t) \leftarrow \text{DIST}_{i-1}(s, t)$ ;
7     foreach  $x \in \text{SEPARATOR}$  do
8        $\text{DIST}_i(s, t) \leftarrow \min\{\text{DIST}_i(s, t), \text{DIST}_{i-1}(s, x) + \text{DIST}_{i-1}(x, t)\}$ 
9 return  $\{(\text{DIST}_{i_{\max}}(s, t))\}_{s, t \in V \times V}$ 
```

Hard

Modified Floyd warshall

Powerful Blackbox

- Again, $h_i = 1.5^i$ is the hop size we target for in phase i
- To make sure our algorithm works, we need to figure out these details
 - How to find separator of size $O(\frac{n \log n}{h_i})$ fast?
 - Is it correct? Will we find a shortest path between a pair of nodes if it has hop $\leq h_i$?


Powerful Blackbox

- How to find Sep of size $O\left(\frac{n \log n}{h_i}\right)$ fast?
- Firstly, we should understand the size and it is straightforward by plugging every shortest path with $\text{hop} \geq \lfloor \frac{h_i}{4} \rfloor$ into Lemma A.1
- In the first phase, we will find our Sep by considering all n^2 pairs, so it will take $O(n^2 h)$ time.
- However, we cannot check all n^2 pairs every time or else the complexity will explode

Powerful Blackbox

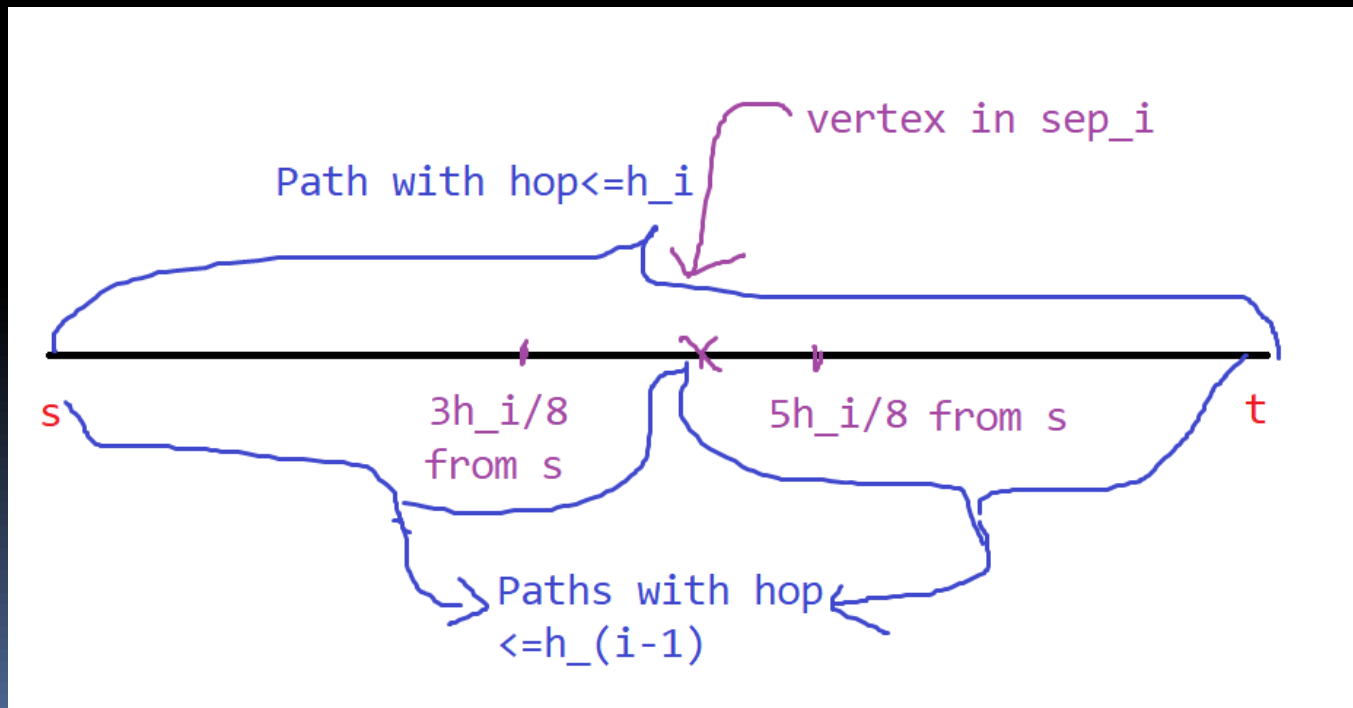
- Lets only check paths that
 - Starts from a vertex in Sep_{i-1} (Sep_{i-1} is the Separator from previous iteration)
 - Ends from a vertex in Sep_{i-1} (Sep_{i-1} is the Separator from previous iteration)
- We take paths that are at least $\frac{h_i}{8}$ hop
- Separator size remains $O\left(\frac{n \log n}{h_i}\right)$
- Time complexity = $O(\log n)$ phases \times
 $O\left(\frac{n \log n}{h_i} n h_i\right) = O(n^2 \log^2 n)$

Powerful Blackbox



Powerful Blackbox

- Is it correct? Will we find a shortest path between a pair of nodes if it has $\text{hop} \leq h_i$?
- Same idea as the one in Delete()



Powerful Blackbox

- Lets analyze complexity
- $O(n^2 h)$ for computing first separator
- $O(n^2 \log^2 n)$ for computing the later separators
- $O\left(\frac{n \log n}{h_i} \times n^2\right) = O\left(\frac{n^3 \log n}{h_i}\right)$ for doing modified floyd warshall
 - Again, analyze with with geometric sum
- $O\left(\frac{n^3 \log n}{h} + n^2 \log^2 n + n^2 h\right)$ total



Further Optimizations

- Optimizes Time Complexity by combining methods with another paper ACK17
 - Optimizes Space Complexity by replacing bellman-ford with an inductive style shortest path (similar as the one in delete), and storing paths in some funny ways
- 